

CM30225 Parallel Computing Shared Memory

Rowan Walshe

November 20, 2017

TODO check valgrind output1

Chapter 1

Code

1.1 Introduction

While working on this coursework, I attempted a number of different solutions for parallelising the problem. In this section I will discuss the different approaches to parallelism for each version, along with the advantages and disadvantages.

1.2 Version 1

In the sequential program of my attempt, the program moves through the 2D array from the top left down to the bottom right. At each point it temporally stores the value at the point, then calculates the new value based on the average of the surrounding four values. It then checks to see if the absolute difference between the new and old value is greater than the given precision. If the value at any point changes more than the given precision, then at the end of this iteration of averaging values, it must complete at least one more iteration. This is illustrated in Code Snippet 1.1.

Code Snippet 1.1: Version 1 Sequential Main Loop

```
while(1) {
    endFlag = true;
    iterations++;
    for(i=1; i<sizeofPlane-1; i++) {
        for(j=1; j<sizeofPlane-1; j++) {
            pVal = plane[i][j];
            plane[i][j] = (plane[i-1][j] + plane[i+1][j]
                + plane[i][j-1] + plane[i][j+1])/4;
            if(endFlag && tolerance < fabs(plane[i][j]-pVal)) {
                endFlag = false;
            }
        }
    }
    if(endFlag) {
        return iterations;
    }
}
```

To parallelise this, each thread is given a certain number of rows that it has calculate. The threads work out which rows of the 2D array they have to work on based on the size of the 2D array, the total

number of threads and their thread ID. For example if the 2D array is of size 10 by 10 and there are four threads total, each thread will work on exactly two rows. If the number of rows does not divide evenly between the number of threads, and there are n rows remaining, then one extra row is given to the thread with ID's less than the number of remaining rows. This means that at most, each thread has to calculate the averages for one extra row. The logic for this can be found in Code Snippet 1.2.

Code Snippet 1.2: Row Split Logic

```

unsigned int  sizeofInner = threadData->sizeofPlane - 2;
unsigned int  rowsPerThreadS = sizeofInner / threadData->threadCount
    + 1;
unsigned int  rowsPerThreadE = sizeofInner / threadData->threadCount;
unsigned int  remainingRows = sizeofInner - threadData->threadCount
    * rowsPerThreadE;

unsigned int  startingRow, endingRow;
if(threadData->id < remainingRows) {
    startingRow = threadData->id * rowsPerThreadS + 1;
    endingRow = startingRow + rowsPerThreadS;
} else {
    startingRow = threadData->id * rowsPerThreadE + remainingRows
        + 1;
    endingRow = startingRow + rowsPerThreadE;
}

```

After a thread has done their portion of the work, they barrier once to wait for all of the threads to finish doing their work on the 2D array. Then each thread checks to see if they need to do another iteration or not. Once they have checked barrier again. The main thread then resets the finished flag, followed by one more barrier. If the threads did not barrier to wait for all the threads to finish their calculations, you end up with a race condition, and some threads may attempt to prematurely return, which would cause the program to freeze. If the threads did not barrier before resetting the finished flag, the flag may be reset before all of the threads have checked whether or not they have finished or not which would be a race condition, and would lead to a similar situation as before. Finally, they barrier once after the flag has been reset, so that the threads do not start work early, which would be a race condition, that may lead to incorrect results. The code for the main loop of the main thread can be found in Code Snippet 1.3. The only difference from the main loop of the child threads is that it counts the number of iterations, and it is the only thread that attempts to reset the finished flag.

Code Snippet 1.3: Version 1 Parallel Main Loop

```

while(1) {
    iterations++;

    relaxPlaneRows(threadData->plane, threadData->sizeofPlane,
        threadData->tolerance, startingRow, endingRow);

    pthread_barrier_wait(&barrierGeneric);
    if(finishedFlag)
        return iterations;
    pthread_barrier_wait(&barrierGeneric);
    finishedFlag = true;
    pthread_barrier_wait(&barrierGeneric);
}

```

While this solution is relatively simple and scales pretty well as can be seen in my scalability testing, it is not particularly fast, with later versions reducing the time it takes to run by upwards of 60%. One of the downsides of this solution is that you may get slightly different results run to run, when using more than one thread. This is because when a thread reaches a row on the boundary between the rows that it has to work on and the rows that another thread is working on, instead of getting update averages above and to the left of it, it now also gets update averages from the row below as well. This means that the sequential and parallel versions of this solution are not really the same algorithm. This means that it may also be harder to prove correctness of the solution.

1.3 Version 2

The goal with my second version was to come up with a way of traversing the 2D array so that no matter the number of threads that the program was using, it would not affect the solution. My idea was to traverse over the 2D array in a chequerboard pattern. This would add some overheads as in the parallel version, you would have to barrier after calculating the average of half of the points in the 2D array, as well as at the end. However, on top of having the added benefit of making correctness testing much easier, it also decreased the time it took to run from my first version by upwards of 50%. The code for the main loop of the sequential program can be found in Snippet 1.4.




Figure 1.1: Chequerboard Pattern

Code Snippet 1.4: Version 2 Sequential Main Loop

```
while(1) {
    finishedFlag = true;
    iterations++;
    for(i=1; i<sizeOfPlane-1; i++) {
        relaxRow(plane, sizeOfPlane, tolerance, i, (i%2)+1);
    }
    for(i=1; i<sizeOfPlane-1; i++) {
        relaxRow(plane, sizeOfPlane, tolerance, i, ((i+1)%2)+1);
    }
    if(finishedFlag) {
        return iterations;
    }
}
```

With this version I created a function *relaxRow* that calculates the new averages for every other value in row *i* starting from the first or second value along in the row depending on which part of the chequerboard is currently being relaxed. The code for the main loop of the parallel version can be found in Snippet 1.5. It bears similarities to the parallel main loop of version one, except now there are four barriers as there is a new barrier half way through the relaxation process.

Code Snippet 1.5: Version 2 Parallel Main Loop

```

while(1) {
    iterations++;
    for(row=startingRow; row<endingRow; row++) {
        relaxPlaneRow(threadData->plane, threadData->sizeOfPlane,
            threadData->tolerance, row, (row%2)+1);
    }
    // Barrier for first half of chequerboard
    pthread_barrier_wait(&barrierGeneric);
    for(row=startingRow; row<endingRow; row++) {
        relaxPlaneRow(threadData->plane, threadData->sizeOfPlane,
            threadData->tolerance, row, ((row+1)%2)+1);
    }
    // Barrier for second half of chequerboard
    pthread_barrier_wait(&barrierGeneric);
    if(finishedFlag) {
        return iterations;
    }
    pthread_barrier_wait(&barrierGeneric);
    finishedFlag = true;
    pthread_barrier_wait(&barrierGeneric);
}

```

As the number of iterations that it has to complete is very similar if not sometimes more than my original version, it is not immediately obvious why this version performs so much better. However by using the cachegrind tool in valgrind, you are able to see that my second version has an L1 data cache miss rate that is less than half that of my first version.

Code Snippet 1.6: Version 1 Cachegrind

```

==65481== Cachegrind, a cache and branch-prediction profiler
==65481== Copyright (C) 2002-2012, and GNU GPL'd, by Nicholas Nethercote et al.
==65481== Using Valgrind-3.8.1 and LibVEX; rerun with -h for copyright info
==65481== Command: ./single -s 300 -u 1 -d 3 -l 4 -r 2 -p 0.00001
==65481==
--65481-- warning: Unknown Intel cache config value (0x63), ignoring
--65481-- warning: L3 cache found, using its data for the LL simulation.
==65481==
==65481== I   refs:      59,145,477,672
==65481== I1 misses:      1,217
==65481== L1i misses:     1,193
==65481== I1 miss rate:    0.00%
==65481== L1i miss rate:  0.00%
==65481==
==65481== D   refs:      27,793,013,948 (24,735,366,146 rd + 3,057,647,802 wr)
==65481== D1 misses:      782,506,106 ( 782,493,703 rd +      12,403 wr)
==65481== L1d misses:     14,256 (      2,163 rd +      12,093 wr)
==65481== D1 miss rate:    2.8% (      3.1% +      0.0% )
==65481== L1d miss rate:  0.0% (      0.0% +      0.0% )
==65481==
==65481== LL refs:      782,507,323 ( 782,494,920 rd +      12,403 wr)
==65481== LL misses:      15,449 (      3,356 rd +      12,093 wr)
==65481== LL miss rate:    0.0% (      0.0% +      0.0% )

```

Code Snippet 1.7: Version 2 Cachegrind

```

==23608== Cachegrind, a cache and branch-prediction profiler

```

```

==23608== Copyright (C) 2002-2012, and GNU GPL'd, by Nicholas Nethercote et al.
==23608== Using Valgrind-3.8.1 and LibVEX; rerun with -h for copyright info
==23608== Command: ./single-chequerboard -s 300 -u 1 -d 3 -l 4 -r 2 -p 0.00001
==23608==
--23608-- warning: Unknown Intel cache config value (0x63), ignoring
--23608-- warning: L3 cache found, using its data for the LL simulation.
==23608==
==23608== I    refs:      59,754,099,652
==23608== I1  misses:      1,215
==23608== L1i misses:      1,192
==23608== I1  miss rate:      0.00%
==23608== L1i miss rate:      0.00%
==23608==
==23608== D    refs:      28,044,004,225 (24,986,267,940 rd + 3,057,736,285 wr)
==23608== D1  misses:      391,254,955 ( 391,242,549 rd +      12,406 wr)
==23608== L1d misses:      14,257 (      2,163 rd +      12,094 wr)
==23608== D1  miss rate:      1.3% (      1.5% +      0.0% )
==23608== L1d miss rate:      0.0% (      0.0% +      0.0% )
==23608==
==23608== LL refs:      391,256,170 ( 391,243,764 rd +      12,406 wr)
==23608== LL misses:      15,449 (      3,355 rd +      12,094 wr)
==23608== LL miss rate:      0.0% (      0.0% +      0.0% )

```

1.4 Version 3

In my final version, I attempted to implement a SWAR using Advanced Vector Extensions (AVX). This would allow me to reduce runtime further by effectively combining up to four arithmetic instructions into one. This added a lot more complexity, as figuring out which parts of the row to use this method on could easily have increased overheads to the point where using AVX would actually become inefficient. In order to keep things simple, I further split up each row into a part that could easily utilise AVX, followed by a parlatenct which would be harder. If the length of the row, ignoring the ends, can be divided by 8 without any remainder, then it is easy to do calculation on 4 value chunks, while still doing a chequerboard pattern. To keep it simple I decided that if there were any remaining points, then I was just fall back to the same method that I used in my second version to calculate the new value for the these points. In practice, any slight slowdown caused by this would become negligible, the larger the problem size. If I had put more time into this, I could have either done SWAR on the remaining points, without utilising the full size of the register, or I could have figured out a way to wrap around onto the next line, continuing to use the full capabilities of AVX as long as possible. The main loop for this version is again very similar, except that two more variables *imax* and *jmax* are passed to the relaxRows function. These tell the function at what point in a row they should stop using AVX and to fall back to the old method. Code Snippets 1.8 and 1.9 show the sequential and parallel main loops respectively.

Code Snippet 1.8: Version 3 Sequential Main Loop

```

remainingItemsPR = sizeofInner%8;
iMax = sizeofPlane-1;
jMax = iMax-remainingItemsPR;
while (1) {
    finishedFlag = true;
    iterations++;
    relaxRows(iMax,jMax,plane,tolerance,0);
    relaxRows(iMax,jMax,plane,tolerance,1);
    if(finishedFlag) {
        return iterations;
    }
}

```

Code Snippet 1.9: Version 3 Parallel Main Loop

```

while (1) {
    iterations++;
    relaxRows(iMax,jMax,threadData->plane,threadData->tolerance,0,
        startingRow,endingRow);
    // Barrier for first half of chequerboard
    pthread_barrier_wait(&barrierGeneric);
    relaxRows(iMax,jMax,threadData->plane,threadData->tolerance,1,
        startingRow,endingRow);
    // Barrier for second half of chequerboard
    pthread_barrier_wait(&barrierGeneric);
    if(finishedFlag) {
        return iterations;
    }
    pthread_barrier_wait(&barrierGeneric);
    finishedFlag = true;
    pthread_barrier_wait(&barrierGeneric);
}

```

As it shows in my scalability testing, this version is not much faster than my second version. However if I was to run my program on an Intel processor with a newer architecture, I would expect it to run faster TODO insert example. Also, If I used a processor that supported AVX-512, I could optimise the program further and do arithmetic on up to 8 packed doubles at once.

Chapter 2

Correctness Testing

I was able to confirm that for small test sizes my final program was correct by manually calculating the result by hand and comparing my result with that of my program. Due to how I am traversing the data in the 2D array, I do not have any race conditions that could affect the result outputted by my program. This means that if the result is correct for the sequential algorithm, then it will also be correct for the parallel algorithm. While I was not able to print out the result of every single run, as I would have run out of storage space, and some of the programs would have run out of time, I did compare a number of results. For each test I also printed out the number of iterations that the program took to reach a solution, so that if there was an error that only happened occupationally, it would have been much easier to spot.

Chapter 3

Scalability Testing